

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#) **FREE TIMED MODE**[Free Practice Exam – Timed Mode](#) [1 Topic](#) **FREE REVIEW MODE**

QUIZ 1 OF 2



FREE AWS Certified Machine Learning Engineer Associate (MLA-C01) (Timed Mode)

9 of 20 questions answered correctly

YOUR TIME:

00:10:51

You have reached 14 of 28 point(s), (50%)

Categories

MLA – Data Preparation for Machine Learning (ML)	66.67%
MLA – Deployment and Orchestration of ML Workflow	25%
MLA – ML Model Development	37.5%

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026


 [Your Progress](#)

 **FREE TIMED MODE**


[Free Practice Exam – Timed Mode](#) [1 Topic](#) 

 **FREE REVIEW MODE**

MLA – ML Solution Monitoring, Maintenance, and Security
ABDULRAH...
Sorry, you failed the test. Carefully read our de explanations including the references and cheat sheets then try again. 🤔



To view your record of all previous attempts:



Visit our FAQ page for more information on the site's features.

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

Restart Quiz

ABDULRAH...

View All Questions

1 ✓ 2 ✓ 3 ✓ 4 ✓ 5 ✓ 6 ✓ 7 ✓ 8 ✓
9 ✓ 10 ✓ 11 ✓ 12 ✓ 13 ✓ 14 ✓ 15 ✓ 16 ✓

✓ Answered  For Review

1. QUESTION

A company is working on a machine learning project that involves analyzing customer data to predict churn rates.

Which of the following statements best describes the machine learning approach the company should take?

☐ Manually analyze the customer data and replicate the analysis process using custom code to automate predictions.

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

☐ Develop a set of hard-coded rules and conditional statements to generate predictions.


☐ Feed the customer data into machine learning algorithms, which will learn patterns and build models to make predictions.

☐ Use raw, unprocessed customer data directly as input to train the machine learning model.

Correct

Machine Learning (ML) is the process of training computers, using math and statistical processes, to find and recognize patterns in data. After patterns are found, ML generates and updates training models to make increasingly accurate predictions and inferences about future outcomes based on historical and new data. For example, ML could help determine the likelihood of a customer purchasing a particular product based on previous purchases by the user or the product's past sales history.

The machine learning process typically involves:



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026


 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#) 

 **FREE REVIEW MODE**

1. **Data collection and preparation:** Gathering relevant data and preprocessing it (cleaning, formatting, feature engineering) to make it suitable for model training.
2. **Model training:** Feeding the prepared data into machine learning algorithms, which learn patterns and relationships within the data to build a model.
3. **Model evaluation:** Assessing the trained model's performance using evaluation metrics and techniques like cross-validation.
4. **Model deployment:** Integrating the trained and evaluated model into applications or systems to make predictions or decisions on new data.



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 FREE TIMED MODE

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 FREE REVIEW MODE

Machine learning involves feeding data into algorithms, which then build models that can make predictions or decisions without being explicitly programmed with rules.

Therefore, the correct answer is: **Feed the customer data into machine learning algorithms, which will learn patterns and build models to make predictions.**

The option that says: **Develop a set of hard-coded rules and conditional statements to generate predictions** is incorrect because machine learning don't just rely on simple hard-coded rules and if-else conditions but rather learns from data and adapts over time. Additionally, AWS emphasizes that machine learning algorithms should not be explicitly programmed with rigid rules.

The option that says: **Manually analyze the customer data and replicate the analysis process using custom code to automate predictions** is incorrect because it only contradicts the automated nature of machine learning. Machine learning algorithms are designed to learn patterns and make predictions automatically from data without the need for manual analysis or replication of human processes.



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 FREE TIMED MODE

[Free Practice Exam – Timed Mode](#) 1 Topic



 FREE REVIEW MODE

The option that says: **Use raw, unprocessed customer data directly as input to the machine learning model** is incorrect because AWS recommends data preprocessing and feature engineering to prepare the data for effective machine learning model training. Using raw, unprocessed data directly can lead to poor model performance and inaccurate predictions.

References:

<https://community.aws/content/2drbbXokwrlXivltJ8ZeCk3gT5F/intro-to-artificial-intelligence-and-machine-learning?lang=en>

<https://aws.amazon.com/what-is/machine-learning/>

Check out this AWS Machine Learning And AI Cheat Sheets:

<https://tutorialsdojo.com/aws-cheat-sheets-aws-machine-learning-and-ai/>

2. QUESTION

A company is developing a machine learning (ML) platform for sentiment analysis of news articles. The platform needs storage to process large volumes of

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 FREE TIMED MODE

[Free Practice Exam – Timed Mode](#) [1 Topic](#)



 FREE REVIEW MODE

text data. The ML platform must handle concurrent access from multiple EC2 instances for data preprocessing, model training, and inference.

ABDUL RAHMAN



Which of the following options will meet the requirements in this scenario?

- ☐ Amazon EBS
- ☐ AWS DataSync
- ☐ Amazon EFS
- ☐ AWS Storage Gateway

Correct

Amazon EFS is a fully managed, scalable, and elastic file storage service that provides seamless integration with existing applications using the Network File System (NFS) protocol. It automatically scales based on usage patterns and replicates data across multiple AZs for high availability.

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)



 **FREE REVIEW MODE**

Amazon Elastic File System

Scalable, elastic, cloud-native NFS file system

Amazon Elastic File System (Amazon EFS) provides a simple, scalable, elastic file system for general purpose workloads for use with AWS Cloud services and on-premises resources.

ABDULRAHMAN

Create file system

Create an EFS file system with the default settings.

Create file system

What is Amazon Elastic File System?



Amazon Elastic File System (Amazon EFS) provides a simple, scalable, fully managed elastic NFS file system for use with AWS Cloud services and on-premises resources. It is built to scale on demand to petabytes without disrupting applications, growing and shrinking automatically as you add and remove files, eliminating the need to provision and manage capacity to accommodate growth. Amazon EFS is designed to provide massively parallel shared access to thousands of Amazon EC2 instances, enabling your applications to achieve high levels of aggregate throughput and IOPS with consistent low latencies.

Pricing

With EFS, there are no minimum fees. You pay only for the storage that you use, the data that you read and write, and any additional throughput that you provision. Estimate your cost using the [AWS Pricing Calculator](#). [Learn more about pricing](#).

Get started

[What is Amazon Elastic File System?](#)
[Get started](#)
[How it works](#)

More resources

[Documentation](#)
[FAQ](#)
[EFS forum](#)
[Contact us](#)

EFS provides high levels of aggregate throughput and IOPS with consistently low latencies by enabling massively parallel shared access to thousands of Amazon EC2 instances for your applications. You can remotely and securely mount EFS file systems to multiple Amazon EC2 instances by using the AWS Systems Manager Run Command.

Hence, the correct answer is: **Amazon EFS**.

The option that says: **Amazon EBS** is incorrect because while EBS volumes are reliable and performant, they are only tied to specific instances and don't support shared access.

The option that says: **AWS Storage Gateway** is incorrect. This service is simply used for connecting on-premises environments to AWS storage services.

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 FREE TIMED MODE

[Free Practice Exam – Timed Mode](#) 1 Topic



 FREE REVIEW MODE

The option that says: **AWS DataSync** is incorrect because it's just a data transfer service for moving data between on-premises and AWS. DataSync is useful for data migration but not designed for concurrent access or ML model training.

References:

<https://docs.aws.amazon.com/efs/latest/ug/whatisefs.html>

<https://aws.amazon.com/efs/faq/>

[https://docs.aws.amazon.com/efs/latest/ug/mount-multiple-ec2-](https://docs.aws.amazon.com/efs/latest/ug/mount-multiple-ec2-instances.html#:~:text=You%20can%20mount%20EFS%20file,AWS%2)

[instances.html#:~:text=You%20can%20mount%20EFS%20file,AWS%2](https://docs.aws.amazon.com/efs/latest/ug/mount-multiple-ec2-instances.html#:~:text=You%20can%20mount%20EFS%20file,AWS%2)

Check out this Amazon EFS Cheat Sheet:

<https://tutorialsdojo.com/amazon-efs/>

3. QUESTION

A data scientist is working for a company that specializes in developing machine learning models for large-scale video processing tasks. The current project involves training a deep-learning model using a massive dataset of high-resolution videos. The dataset is stored in a distributed file system, and the training process requires high throughput and low latency access to the data. Additionally, the training

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

process involves frequent read and write operations and the storage solution must be able to handle high I/O demands efficiently.

ABDUL RAH



Which of the following storage services will meet the requirements?

- ☐ Amazon File Cache
- ☐ Amazon FSx for Lustre
- ☐ AWS Storage Gateway
- ☒ Amazon Elastic File System (Amazon EFS)

Incorrect

Amazon FSx for Lustre is designed specifically for high-performance computing (HPC) and machine learning workloads. It provides sub-millisecond latencies and supports hundreds of gigabytes per second of throughput, which is ideal for training deep learning models with large datasets.



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026


 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)



 **FREE REVIEW MODE**



FSx for Lustre efficiently handles the high I/O operations required for processing high-resolution videos. It is optimized for both sequential and random access workloads. It can seamlessly integrate with Amazon S3, allowing for efficient data access and management, which is beneficial when dealing with massive datasets typically used in machine learning.

Hence, the correct answer is: **Amazon FSx for Lustre.**

The option that says: **Amazon Elastic File System (Amazon EFS)** is incorrect. While Amazon EFS provides scalable file storage for use with AWS Cloud services and on-premises resources, it is designed more for general-purpose workloads rather than high-performance computing.

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#) FREE TIMED MODE

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 FREE REVIEW MODE

The option that says: **AWS Storage Gateway** is incorrect because it is primarily designed to connect on-premises environments to AWS storage services, providing seamless integration and hybrid cloud storage solutions. It is not optimized for high-performance computing or the specific needs of deep learning model training. Its primary use case is for data backup, archiving, and hybrid cloud storage scenarios.

The option that says: **Amazon File Cache** is incorrect because it is only used to cache frequently accessed data, which does not align with the requirements of continuous and intensive read/write operations needed for training deep learning models.

References:

<https://docs.aws.amazon.com/fsx/latest/LustreGuide/what-is.html>

<https://docs.aws.amazon.com/efs/latest/ug/whatisefs.html>

Check out this Amazon FSx Cheat Sheet:

<https://tutorialsdojo.com/amazon-fsx/>

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#) **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

A data scientist is in the process of examining a machine learning model to identify potential bias within the training data before it's deployed in a production environment. The goal is to guarantee that the model exhibits no disproportionate favoritism or discrimination towards any specific demographic group.

Match each pre-training bias metric with its appropriate scenario. Each metric may be used more than once. (Select THREE.)

- Total Variation Distance (TVD)
- Kullback-Leibler Divergence (KL)
- Difference in Proportions of Labels (DPL)
- Class Imbalance (CI)

The data scientist wants to compare the difference in the proportion of positive outcomes among different racial groups.

Difference in Proportions of Labels (DPL)

The dataset must be analyzed to measure the imbalance between male and female instances.

Class Imbalance (CI)

The data scientist is interested in evaluating the maximum disparity in results among different income

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 Your Progress

 FREE TIMED MODE

[Free Practice Exam – Timed Mode](#) 1 Topic

 FREE REVIEW MODE

levels.


ABDULRAH...
Kullback-Leibler Divergence (KL)

Variation Distance (TVD))

Incorrect

2 / 3 Points

Evaluating and reducing bias in machine learning models is crucial for creating fair and accurate models, especially when dealing with sensitive demographic data. Pre-training bias metrics are essential tools for data scientists to identify and address potential biases within a dataset before the model is trained. Applying these metrics can help ensure that machine learning models are more equitable and that predictions do not disproportionately favor or disadvantage specific groups.



The screenshot shows the AWS Data Wrangler interface. The main panel displays a table with shipping data. The 'Create analysis' sidebar on the right lists several bias metrics:

- ☒ Class imbalance (CI)
- ☒ Difference in Positive Proportions in Labels (DPL)
- ☐ IS divergence (IS)
- ☐ Conditional Demographic Disparity in Labels (CDDL)
- ☐ Would you like to analyze additional metrics?
 - ☒ Yes
 - ☐ No
- ☐ Kullback-Leibler Divergence (KL)
- ☐ Lp-norm (LP)
- ☒ Total Variation Distance (TVD)
- ☐ Kolmogorov-Smirnov Distance (KS)

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

Class Imbalance (CI) – This metric evaluates the representation of different classes or demographic groups in a dataset. For instance, if one gender is significantly underrepresented, CI will reveal this imbalance, which could lead to biased predictions that favor the more represented class. Addressing such imbalances is important to ensure that the model learns fairly from all data groups.

Difference in Proportions of Labels (DPL) – DPL measures disparities in the proportion of positive outcomes between different demographic groups. For example, if a model predicts higher success rates for one racial group compared to another, DPL will highlight this discrepancy, indicating a potential bias in outcomes. Analyzing DPL can make adjustments to the model or dataset to reduce such disparities and promote fairness.

Total Variation Distance (TVD) – TVD quantifies the maximum divergence in outcome distributions between different demographic groups. It's a critical metric for identifying how much the predicted outcomes for one group differ from another, thus helping to assess bias in the model's predictions.

Hence, the correct answers are:



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

– The data scientist wants to compare the difference in the proportion of outcomes among different racial groups: **Difference in Proportions of Labels (DPL)**

– The dataset must be analyzed to measure the imbalance between male and female instances: **Class Imbalance (CI)**

– The data scientist is interested in evaluating the maximum disparity in results among different income levels: **Total Variation Distance (TVD)**

The option that says: **Kullback-Leibler Divergence (KL)** is incorrect. This metric primarily measures the divergence between two probability distributions often used to compare the distribution of outcomes across different groups. KL Divergence is not specifically tailored for measuring class imbalance or comparing the proportion of labels across demographic groups. It is typically used in broader contexts of distribution comparisons and might not provide a direct assessment of bias related to demographic attributes.

References:



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

<https://docs.aws.amazon.com/sagemaker/latest/measure-data-bias.html> [ABDULRAH...](#)

<https://docs.aws.amazon.com/sagemaker/latest/detect-data-bias.html>

Check out this AWS Machine Learning And AI
Cheat Sheets:

<https://tutorialsdojo.com/aws-cheat-sheets-aws-machine-learning-and-ai/>

5. QUESTION

A retail company is launching a series of new products across its global stores. The company wants to forecast daily demand for its products based on historical sales data. The historical data includes information on product sales, inventory levels, and marketing campaigns. However, the timestamps for this data are irregular, and some data points are missing due to occasional stockouts or incomplete records.

Which approach would you recommend to satisfy these requirements with the LEAST operational overhead?



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

☐ Visualize and forecast data using Amazon QuickSight. [ABDULRAH...](#)

☐ Clean and forecast data using Amazon SageMaker Studio Notebook with Pandas.

☐ Process and forecast data using Amazon EMR Serverless with PySpark.

☐ Prepare and analyze data using Amazon SageMaker Studio Data Wrangler.

Correct

Amazon SageMaker Studio Data Wrangler provides a streamlined and visual interface for data preparation and analysis. It enables users to handle irregular timestamps and missing data points effectively through built-in transformations and data cleaning functions. Data Wrangler is designed to minimize operational overhead by offering a low-code solution for data wrangling tasks, allowing users to quickly prepare data for machine learning without deep technical expertise in data engineering or programming.

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete


Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**



By leveraging Data Wrangler, the corporation can ensure efficient data preparation and focus more on the analysis and forecasting aspects with reduced operational overhead.

Hence, the correct answer is: **Prepare and analyze data using Amazon SageMaker Studio Data Wrangler.**

The option that says: **Process and forecast data using Amazon EMR Serverless with PySpark** is incorrect because while this can primarily handle large-scale data processing and complex transformations, it requires more setup and maintenance compared to Data Wrangler. This approach involves configuring clusters, managing PySpark scripts, and ensuring infrastructure scalability, leading to higher operational overhead.

The option that says: **Visualize and forecast data using Amazon QuickSight.** is incorrect. Amazon QuickSight is just a tool for data visualization with basic forecasting capabilities. It is not primarily

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#) **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

designed for extensive data cleaning or handling irregular timestamps and missing data points.

QuickSight focuses more on presenting insights rather than preparing data, which is a crucial step in this scenario.

The option that says: **Clean and forecast data using Amazon SageMaker Studio Notebook with Pandas** is incorrect. This option requires significant coding and manual intervention to handle irregular timestamps and missing data points compared to the streamlined interface provided by Data Wrangler.


References:

<https://docs.aws.amazon.com/sagemaker/latest/dg/data-wrangler.html>

<https://aws.amazon.com/blogs/machine-learning/prepare-time-series-data-with-amazon-sagemaker-data-wrangler/>

Check out this Amazon SageMaker Cheat Sheet:

<https://tutorialsdojo.com/amazon-sagemaker/>



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)


 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)



 **FREE REVIEW MODE**

A Machine Learning Specialist uses Amazon SageMaker Data Wrangler to prepare the data to train a model. The objective is to analyze the data, understand feature relationships and detect issues. The goal is to visualize correlations and assess strength to identify patterns and outliers.



ABDUL RAH

Which visualization technique should be used?

- ☐ Scatter Plot
- ☐ Histogram
- ☐ Multicollinearity
- ☐ Shapley values plot

Correct

Amazon SageMaker Data Wrangler offers a visual interface for preparing data for machine learning. It allows users to perform various data analysis tasks, including data visualization. Visualizing relationships between variables and evaluating the strength of these relationships is essential for understanding the dataset and spotting patterns and outliers. A **scatter plot** is a useful tool for

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026


 Your Progress

 FREE TIMED MODE

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 FREE REVIEW MODE

showing the relationship between two continuous variables. By graphing each observation as a point, the scatter plot helps in identifying correlations, trends, and potential outliers in the data. It provides a clear visual representation of how one variable changes in relation to another, making it easier to detect linear and non-linear relationships, as well as clusters and deviations.



In SageMaker Data Wrangler, scatter plots help understand feature relationships for feature engineering and model selection. Strong correlations may indicate redundant variables, while weak correlations may signal the need for additional features or data transformations. Additionally, scatter plots can identify outliers for effective dataset cleaning before model training, leading to more accurate predictive models.

Hence, the correct answer is: **Scatter Plot.**

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

Histogram is incorrect. This option is only used to visualize the distribution of a single variable rather than relationships between two variables. While histograms are helpful for understanding the distribution and frequency of data points, they do not provide insight into how two variables interact or correlate with each other.

Shapley values plot is incorrect. While Shapley values are highly effective for understanding feature importance and attribution after a model is trained (typically via SageMaker Clarify), they are not a standard exploratory data visualization technique in Data Wrangler for initially assessing raw feature correlations.

Multicollinearity is incorrect because it's just a statistical phenomenon in which several independent variables in a model are highly correlated. It is not a visualization technique but rather an issue that needs to be detected and addressed through statistical analysis, often using techniques like variance inflation factor (VIF) or correlation matrices.

References:

<https://docs.aws.amazon.com/sagemaker/latest/dg/data-wrangler-analyses.html#data-wrangler-visualize-scatter-plot>



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#) **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

<https://docs.aws.amazon.com/sagemaker/latest/>

[ABDULRAH...](#)

**Tutorials Dojo's AWS Machine Learning Cheat
Sheets:**

<https://tutorialsdojo.com/aws-cheat-sheets-aws-machine-learning-and-ai/>

7. QUESTION

A manufacturing company has uploaded raw event log data into a data repository to build a machine learning (ML) model for predictive maintenance. The raw event logs include vehicle numbers, scheduled maintenance, warranty expiration, wheel rotation, and more. The goal is to transform the data into an efficient format for storing and processing.

What data format should be used?

☐ Parquet

☐ XML

☐ CSV Format

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#) **FREE TIMED MODE**[Free Practice Exam – Timed Mode](#) [1 Topic](#) **FREE REVIEW MODE**

Incorrect

Parquet is a columnar storage file format optimized for big data processing frameworks such as Apache Hadoop, Apache Spark, and Apache Drill. It offers excellent data compression and encoding algorithms, which help reduce storage space and improve query performance.

You can use AWS Glue Parquet writer to enable faster Parquet file writing. The conventional writer creates a schema before writing. The Parquet format does not keep the schema in an easily retrievable format; therefore, this may take some time. The AWS Glue Parquet writer does not require a pre-computed schema. As new data arrives, the writer computes and updates the schema dynamically.



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete


Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**



```
1 from pyspark.context import SparkContext
2 from awsglue.context import GlueContext
3
4 sc = SparkContext.getOrCreate()
5 glueContext = GlueContext(sc)
6 spark = glueContext.spark_session
7
8 dynamicFrame = glueContext.create_dynamic_frame_from_options(
9     connection_type = "s3",
10    connection_options = {"paths": ["s3://s3path/"]},
11    format = "parquet"
12 )
```

Parquet's columnar format is ideal for storing and processing large datasets efficiently. It allows for high compression and faster query performance, making it suitable for ML models where fast access to specific data columns is crucial. This efficiency aligns to transform raw event logs into an optimal format for predictive maintenance.

Hence, the correct answer is: **Parquet**.

The option that says: **CSV Format** is incorrect. While it is easy to use and widely supported, it is not as efficient as columnar formats for large-scale data processing and storage. It does not support advanced data compression and indexing, which can lead to larger storage requirements and slower query performance.

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#) FREE TIMED MODE

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 FREE REVIEW MODE

The option that says: **XML** is incorrect because XML is primarily a hierarchical **ABDULRAHMAN** format that is verbose and not optimized for storage and processing efficiency, especially with large datasets. While XML is good for data interchange and readability, it does not support efficient data compression and query performance, making it less suitable for big data processing and analytics tasks compared to columnar storage formats like Parquet.

The option that says: **Avro RecordIO** is incorrect because this option is only a row-based storage format that is efficient for serialization but not ideal for storage and processing of large datasets like Parquet. Avro does support schema evolution, but its row-based nature makes it less efficient for large-scale analytical workloads compared to Parquet.

References:

<https://docs.aws.amazon.com/glue/latest/dg/aws-glue-programming-etl-format-parquet-home.html>

<https://parquet.apache.org/docs/overview/>

Check out this **AWS Glue DataBrew Cheat Sheet:**

<https://tutorialsdojo.com/aws-glue-databrew/>

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

8. QUESTION

[ABDULRAH...](#)



An e-commerce company wants to implement a machine learning system to detect fraudulent online transactions. An ML engineer has deployed a machine learning model to identify and minimize these fraudulent activities. However, only a small percentage (around 3%) of the transactions in the dataset are found fraudulent. The company wants to fine-tune the model's ability to detect fraud while minimizing false positives to ensure genuine transactions are not impeded.

Which metrics should the ML engineer use to optimize the performance of its fraud detection system? (Select TWO.)

☐ True Positive Rate

☐ F1 Score

☒ Cut-off

☒ RMSE (Root Mean Square Error)

☐ False Positive Rate

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

Incorrect

ABDULRAH...

RMSE (Root Mean Square Error) is a metric used to measure the differences between predicted values and actual values in a regression problem. It calculates the square root of the average squared differences between the predicted and actual values. A lower Root Mean Square Error value indicates better model performance.

True Positive Rate (Recall or Sensitivity), also known as Recall or Sensitivity, is a metric used in classification problems. It assesses the percentage of true positive instances correctly identified by the model. A higher True Positive Rate indicates that the model is better at detecting positive instances.

False Positive Rate is another metric used in classification problems. It measures the proportion of actual negative instances that are incorrectly classified as positive by the model. A lower False Positive Rate is desirable, as it means the model is less likely to misclassify negative instances as positive.

F1 Score is a metric that merges recall and precision into a single score. It is the harmonic mean of recall and precision. Precision measures the fraction of correctly predicted positive outcomes, while recall is the same as the True

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 Your Progress


 FREE TIMED MODE

[Free Practice Exam – Timed Mode](#) 1 Topic

 FREE REVIEW MODE

Positive Rate. The F1 Score provides a balanced measure of a model's performance by considering both precision and recall.

Cut-off (or threshold) is a value used to convert the model's predicted probabilities or scores into binary class predictions. Specifically, the cut-off is a threshold value that separates the two classes. Any predicted probability or score above the cut-off value is classified as one class (e.g., positive or fraudulent), while any value below the cut-off is classified as the other class (e.g., negative or genuine).



You can provide your training data as a file or pipe.

File

Metrics

The algorithm you selected will publish the following metrics to CloudWatch metrics.

Metric name	Regex
test:f1	#quality_metric: host=[\$+], test f1 <score>=([\$+])
train:progress	#progress_metric: host=[\$+], completed (\$+) %
train:recall	#quality_metric: host=[\$+], train recall <score>=([\$+])

☐ Enable SageMaker metrics time series

Resource configuration

Instance type: ml.m4.xlarge, Instance count: 1, Additional storage volume per instance (GB): 1

Keep alive period: 1 minutes

Maximum runtime is 1 hour (60 minutes or 3600 seconds).

Encryption key - optional

Encrypt your data. Choose an existing KMS key or enter a key's ARN.

No Custom Encryption

Stopping condition

Maximum runtime: 24 hours

The problem statement mentions that the company wants to fine-tune the fraud detection model to identify fraudulent transactions while minimizing false positives (genuine transactions incorrectly flagged as fraudulent). By optimizing (minimizing)

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**


the FPR, the ML engineer can ensure that the fraud detection system does not unnecessarily flag many genuine transactions as fraudulent. Remember that FPR directly measures the proportion of genuine (non-fraudulent) transactions that are incorrectly classified as fraudulent by the model. A low FPR indicates that the model is minimizing the number of false positives, which aligns with the company's requirement. Next is optimizing the F1 Score which helps strike a balance between Precision (minimizing false positives) and Recall (detecting true fraudulent cases). By maximizing the F1 Score, the ML engineer can optimize the model to have both a reasonably high Precision (low false positives) and a high Recall (high true positive rate for fraudulent transactions).

Hence, the correct answers are:

– False Positive Rate

– F1 Score

The option that says: **RMSE (Root Mean Square Error)** is incorrect because it is a metric primarily used for regression problems, where the goal is to predict a continuous value. In the case of fraud detection, which is a classification problem, RMSE is not an appropriate metric.



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

The option that says: **True Positive Rate** is incorrect. The True Positive Rate (TPR) measures how well the model identifies actual fraudulent transactions but does not address the need to minimize false positives. A high TPR indicates that the model detects many fraudulent cases accurately; however, it can also result in a high False Positive Rate (FPR), meaning that genuine transactions might be wrongly flagged as fraudulent. Since the company's goal is to minimize false positives to avoid disrupting legitimate transactions, focusing on metrics like FPR and F1 Score is more appropriate. These metrics directly address the issue of false positives and ensure that the model effectively balances detecting fraud while minimizing incorrect flags on genuine transactions.

The option that says: **Cut-off** is incorrect. The cut-off (or threshold) is not a performance metric but only a decision boundary used to classify transactions as fraudulent or not. Adjusting the cut-off is part of the process to optimize performance metrics like precision, recall, and F1 Score, but it is not a metric itself.

References:

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

<https://docs.aws.amazon.com/machine-learning/latest/dg/amazon-machine-learning-concepts.html>

<https://docs.aws.amazon.com/machine-learning/latest/dg/what-is-amazon-machine-learning.html>

Check out this AWS Machine Learning and AI Cheat Sheet:

<https://tutorialsdojo.com/aws-cheat-sheets-aws-machine-learning-and-ai/>



9. QUESTION

A company wants to use cost allocation tags to manage and optimize its AWS costs related to Machine Learning (ML) workloads more effectively. It also aims to track and report on its spending across different ML projects and departments.

Select and order the correct steps from the following list to implement cost allocation tagging for this purpose. Each step should be selected one time or not at all. (Select and order THREE.)

- Create user-defined tags (key-value pairs) for the ML resources.

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

- Configure tag-based cost and usage report detailed analysis. [ABDULRAH...](#)
- Enable the cost allocation tags in the AWS E console.
- Set up automated alerts based on budget thresholds for ML expenses.
- Attach the tags to the ML resources in the AWS Management Console.

1. **Attach the tags to the ML resources in the AWS M**

(Create user-defined tags (key-value pairs) for the ML resources.)

2. **Configure tag-based cost and usage reports for d**

(Attach the tags to the ML resources in the AWS Management Console.)

3. **Enable the cost allocation tags in the AWS Billing**

Incorrect

1 / 3 Points

AWS Cost Management is a crucial service for organizations using AWS, as it helps them gain visibility into their cloud spending, identify cost-saving opportunities, and make informed decisions

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026


 Your Progress

 FREE TIMED MODE

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 FREE REVIEW MODE

about resource allocation and optimization. It provides various tools and reports, such as Cost Explorer, Cost and Usage Reports, and Budgets, help users analyze and manage their AWS costs effectively. projects or departments. Without proper cost management and governance, your ML spend might lead to surprises on your monthly AWS bill.



Cost allocation tagging is a feature in AWS that allows you to assign metadata (tags) to your AWS resources. These tags can then be used to track and report on costs associated with those resources. This is particularly useful for organizations that need to allocate costs across different projects, departments, or cost centers. By tagging your resources, you can easily identify and report on the costs associated with specific projects, departments, or environments.

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 FREE TIMED MODE

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 FREE REVIEW MODE

To implement cost allocation tags, the initial steps are as follows: First, create user-defined tags, which are key-value pairs that reflect attributes such as project names or departments to ensure proper categorization of ML resources. Then, apply these tags to the relevant ML resources to enable tracking. Finally, enable the cost allocation tags in the Billing console, which ensures that the tags are included in cost allocation and reporting, facilitating accurate tracking and analysis of costs associated with ML workloads.

Hence, the correct answers are:

- Step 1: **Create user-defined tags (key-value pairs) for the ML resources.**
- Step 2: **Attach the tags to the ML resources in the AWS Management Console.**
- Step 3: **Enable the cost allocation tags in the AWS Billing console.**

The option that says: **Configure tag-based cost and usage reports for detailed analysis** is incorrect. This step should come after the first three steps. You need to enable cost allocation tags, create the tags, and attach them to the ML resources before you can use Cost Explorer to analyze and report costs based on those tags.



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#) **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

The option that says: **Set up automated alerts based on budget thresholds for AWS expenses** is incorrect. While important for managing costs, this option is not a necessary step for the tagging process. It's more related to budget management and monitoring rather than cost allocation tagging.

References:

<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/custom-tags.html>

<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/cost-alloc-tags.html>

Check out this AWS Billing and Cost Management Cheat Sheet:

<https://tutorialsdojo.com/aws-billing-and-cost-management/>

10. QUESTION

A financial adviser plans to use machine learning to improve the efficiency and accuracy of processing mortgage loan requests. By incorporating relevant financial and behavioral data, they aim to reduce the time needed to evaluate applications.

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#) 

 **FREE REVIEW MODE**

Which options are important for gathering data for model training?

[ABDULRAH...](#)



☐ Education History

☐ Medical history

☐ Credit bureau data

☐ Employment history

Correct

Credit bureau data is crucial in evaluating an individual's financial behavior and risk profile. It includes information about credit scores, outstanding debts, and payment histories. In machine learning, the quality and relevance of the data used to train the model are critical for achieving accurate predictions. For loan application processing, financial and behavioral data, including credit bureau data, provide insights into an applicant's ability to repay a loan.



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete


Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**



Credit bureau data is essential for creating accurate risk assessment models. It provides detailed insights into an individual's financial stability and creditworthiness, making it highly relevant for training machine learning models in loan processing. This data helps in predicting the likelihood of defaults, thereby improving the accuracy of risk assessments and decision-making processes.

Hence, the correct answer is: **Credit bureau data.**

The option that says: **Employment history** is incorrect. While employment history provides some information about job stability and income, it does not offer as comprehensive a picture of financial behavior as credit bureau data. It is less directly relevant to financial risk assessment.

The option that says: **Medical history** is incorrect because this is related to an individual's health and has no direct correlation with their financial

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

behavior or creditworthiness. In addition, it would likely violate privacy regulations. **ABDULBAHIS** data

The primary focus here is on financial and behavioral data that directly impact the likelihood of defaults or claims rather than health-related data.

The option that says: **Education History** is incorrect. While it offers insights into earning potential and stability, it is not as directly relevant to assessing financial behavior and creditworthiness as other factors like credit bureau data, which provide crucial indicators of financial status and repayment ability for mortgage loan processing.

References:

<https://www.experian.com/blogs/insights/what-is-credit-risk-analytics/>

<https://aws.amazon.com/blogs/machine-learning/prepare-data-for-predicting-credit-risk-using-amazon-sagemaker-data-wrangler-and-amazon-sagemaker-clarify/>

Check out this Amazon SageMaker Cheat Sheet:

<https://tutorialsdojo.com/amazon-sagemaker/>



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)


 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

11. QUESTION

[ABDULRAH...](#)



A machine learning engineer is working on an AWS project involving real-time fraud detection for an online payment platform. The platform generates a continuous stream of transaction data that needs processing and analysis using a pre-trained model on Amazon SageMaker AI.

Which combination of AWS services should the machine learning engineer choose to ingest and process the real-time transaction data before delivering it to SageMaker AI for inference? (Select TWO.)

- ☐ Amazon Data Firehose
- ☐ AWS Glue
- ☒ Amazon Managed Streaming for Apache Kafka (Amazon MSK)
- ☐ Amazon Kinesis Data Streams
- ☐ Amazon Managed Service for Apache Flink

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

Incorrect

ABDULRAH...

Amazon Kinesis Data Streams is a scalable and durable real-time data streaming service offered by AWS. It allows you to continuously capture and store terabytes of data per hour from hundreds of thousands of sources, such as website clickstreams, database event streams, financial transactions, social media feeds, and IoT sensor data.

Kinesis Data Streams can automatically scale to handle varying data volumes and provide features like data retention, data replication, and automatic load balancing. It is designed to handle real-time data ingestion from various sources and can be integrated with other AWS services like Amazon Managed Service for Apache Flink Studio, AWS Lambda, and Amazon SageMaker for real-time data processing and analysis.

Apache Flink is an open-source distributed processing engine for stateful computations over data streams. It provides a high-performance runtime and a powerful stream processing API that supports stateful computations, event-time processing, and accurate fault-tolerance guarantees.

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 Your Progress

 FREE TIMED MODE


Free Practice Exam – Timed Mode 1 Topic

 FREE REVIEW MODE

Amazon Managed Service for Apache Flink is

fully managed service that enables you to run Apache Flink applications on AWS without having to manage the underlying infrastructure. It automatically provisions and manages the resources required to run Apache Flink clusters, including EC2 instances, Apache Flink runtime, and Apache ZooKeeper.

With Amazon Managed Service for Apache Flink, you can consume data from various sources, including Amazon Kinesis Data Streams, and perform real-time data transformations, filtering, and enrichment. It supports stateful stream processing, allowing you to maintain and update the state as data flows through the application. This service integrates with other AWS services like Amazon SageMaker, Amazon S3, and Amazon Data Firehose, enabling you to build end-to-end real-time data processing pipelines.



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

By using Kinesis Data Streams and Amazon Managed Service for Apache Flink together, the machine learning engineer can ingest the real-time transaction data from the online payment platform, apply any necessary transformations or preprocessing, and then deliver the processed data to Amazon SageMaker for real-time inference using the pre-trained machine learning model to detect fraudulent transactions.


Hence, the correct answers are:

- Amazon Kinesis Data Streams
- Amazon Managed Service for Apache Flink

Amazon Data Firehose is incorrect. While Kinesis Data Firehose can ingest and deliver data to Amazon S3 or other destinations, it is primarily designed for batch processing and not real-time inference.

AWS Glue is incorrect because AWS Glue is a fully managed extract, transform, and load (ETL) service, primarily used for batch data processing and not real-time data ingestion and processing.

Amazon Managed Streaming for Apache Kafka (Amazon MSK) is incorrect. While this service could technically handle streaming data, it's mainly



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#) **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

used for data transfer between different applications and systems rather than real-time processing and inference.

References:

<https://docs.aws.amazon.com/streams/latest/dev/introduction.html>

<https://docs.aws.amazon.com/whitepapers/latest/big-data-analytics-options/amazon-kinesis.html>

<https://aws.amazon.com/managed-service-apache-flink/?nc=sn&loc=1>

Check out these AWS Cheat Sheets:

<https://tutorialsdojo.com/amazon-kinesis/>

<https://tutorialsdojo.com/amazon-kinesis-data-streams-vs-data-firehose-vs-data-analytics-vs-video-streams/>

12. QUESTION

A large retail company is exploring generative AI to enhance its customer service chatbot. The company aims to fine-tune models from third-party services to enhance the chatbot's comprehension of business-specific customer inquiries.

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

Which benefits of Amazon Bedrock will help the company meet its requirements? (Select TWO).

☐

Offers distinct APIs for accessing different foundation models (FMs).

☐

Amazon Bedrock enables private customization of foundation models (FMs).

☐

Ensures that prompts and responses remain completely private and secure.

☒

Amazon Bedrock requires customers to handle infrastructure provisioning.

☐

Offers a wide range of pre-trained foundation models (FMs) from multiple providers.

Incorrect

Amazon Bedrock is a service that provides access to foundation models (FMs) from leading AI providers. It allows businesses to integrate and scale generative AI applications without the complexity of managing underlying infrastructure.

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

With Amazon Bedrock, companies can easily access various pre-trained models such as text generation, chatbots, image recognition, and more. This service offers flexibility in choosing the best models for specific use cases, enabling organizations to leverage cutting-edge AI technology to enhance their applications and customer experiences.



One of the primary benefits of Amazon Bedrock is its support for the private customization of foundation models. This means businesses can fine-tune pre-trained models to suit their unique needs and data better, ensuring more accurate and relevant outputs. Additionally, Amazon Bedrock ensures data privacy and security, as prompts and responses are handled securely within the AWS environment. This allows companies to deploy AI solutions while confidently complying with data protection regulations. Amazon Bedrock empowers

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 FREE TIMED MODE

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 FREE REVIEW MODE

businesses to innovate rapidly and effectively in the AI landscape by providing a broad selection of foundation models and robust security features.


Hence, the correct answers are:

- Offers a wide range of pre-trained foundation models (FMs) from multiple providers.
- Amazon Bedrock enables private customization of foundation models (FMs).

The option that says: **Ensures that prompts and responses remain completely private and secure** is incorrect. While important, the scenario does not specifically highlight privacy as a concern for enhancing the chatbot's functionality. Therefore, it's not among the most directly relevant benefits for the given requirement.

The option that says: **Offers distinct APIs for accessing different foundation models (FMs)** is incorrect. This option simply describes a way of getting details about different foundation models available on Amazon Bedrock.

The option that says: **Amazon Bedrock requires customers to handle infrastructure provisioning** is incorrect. Amazon Bedrock is primarily designed to take care of the infrastructure provisioning and



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete


Last activity on September 5, 2026

 [Your Progress](#) **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

management. Customers do not need to handle underlying infrastructure, allowing them to focus on developing and deploying their AI applications.



References:

<https://docs.aws.amazon.com/bedrock/latest/userguide/what-is-bedrock.html>

<https://aws.amazon.com/bedrock/>

Check out this Amazon Bedrock Cheat Sheet:

<https://tutorialsdojo.com/amazon-bedrock/>

13. QUESTION

A healthcare company seeks to enhance patient outcome predictions using generative AI applications. The company requires a solution that allows selection from various predictive models, guarantees the confidentiality of private data during model fine-tuning, and eliminates the need for managing the underlying ML infrastructure.

Which AWS service best meets the requirements?

☐ Amazon Comprehend Medical

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) 1 Topic

 **FREE REVIEW MODE**

☐ Amazon Rekognition


[ABDULRAH...](#)

☐ Amazon Bedrock

☐ Amazon SageMaker Studio

Correct

Amazon Bedrock is a fully managed service that offers leading foundation models (FMs) and a set of capabilities to quickly build and scale generative artificial intelligence (generative AI) applications.



Amazon Bedrock
The easiest way to build and scale generative AI applications with foundation models (FMs)

Try Bedrock
[Get started](#)

Overview

Amazon Bedrock is a fully managed service that makes FMs from leading AI startups and Amazon available via an API, so you can choose from a wide range of FMs to find the model that is best suited for your use case. With Bedrock's serverless experience, you can get started quickly, privately customize FMs with your own data, and easily integrate and deploy them into your applications using the AWS tools without having to manage any infrastructure.

Benefits

- Accelerate development of generative AI applications using FMs through an API, without managing infrastructure.
- Choose FMs from AI21 Labs, Anthropic, Stability AI, and Amazon to find the right FM for your use case.

Amazon Bedrock provides a serverless environment, simplifying the deployment and management of machine learning models without the need to handle the underlying infrastructure. Bedrock supports customizations of these models,

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**


ensuring privacy and security for user data, making it suitable for various applications, including those requiring handling of sensitive information.

Hence, the correct answer is **Amazon Bedrock**.

The option that says: **Amazon SageMaker Studio** is incorrect. While this service offers comprehensive tools for machine learning, it's primarily designed for broad ML tasks rather than focusing specifically on generative AI and foundation models. It requires more hands-on management of the machine learning lifecycle compared to Bedrock.

The option that says: **Amazon Comprehend Medical** is incorrect because this service is tailored only for extracting medical information from unstructured text using natural language processing. It does not support the creation or management of generative AI models, as it's focused solely on NLP tasks in the medical field.

The option that says: **Amazon Rekognition** is incorrect because Amazon Rekognition is simply used for tasks like facial analysis, object detection, and activity recognition. It does not facilitate the building or scaling of foundation models for generative AI applications, making it unsuitable for predictive analytics in healthcare beyond visual data.



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)



 **FREE REVIEW MODE**

References:

<https://docs.aws.amazon.com/bedrock/latest/userguide/what-is-bedrock.html>

<https://aws.amazon.com/bedrock/>

Check out this Amazon Bedrock Cheat Sheet:

<https://tutorialsdojo.com/amazon-bedrock/>



at-

14. QUESTION


A financial services company is leveraging AWS machine learning services to enhance its operations and customer experience. The company is focused on implementing advanced solutions to detect fraudulent activities, gain insights from customer feedback, and offer personalized financial advice through an interactive chatbot.

Select the correct AWS service to meet this requirement for each case. Each service should be selected one or more times. (Select THREE.)

– Amazon SageMaker AI

– Amazon Comprehend

– Amazon Lex



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#) **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

Detecting fraudulent transactions **ABDULRAH...**

Amazon SageMaker AI

Analyzing customer feedback for sentiment and

insights **Amazon Lex** (Amazon

Comprehend)


Providing personalized financial advice through a

chatbot **Amazon Comprehend** (Amazon Lex)

Incorrect

1 / 3 Points

Amazon SageMaker AI is designed to build, train, and deploy machine learning models. It is well-suited for tasks like fraud detection, where complex models can analyze transaction patterns and identify anomalies.



Amazon Comprehend is an NLP (natural language processing) service that analyzes text to extract

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026


 Your Progress

 FREE TIMED MODE


[Free Practice Exam – Timed Mode](#) 1 Topic

 FREE REVIEW MODE

insights like sentiment, key phrases, and entities. It's perfect for evaluating customer feedback.



Amazon Lex is used to build conversational interfaces and chatbots. This can be integrated with other AWS services to provide personalized responses and advice to users.



Hence, the correct answers are:

- Amazon SageMaker AI – Detecting fraudulent transactions
- Amazon Comprehend – Analyzing customer feedback for sentiment and insights

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#) FREE TIMED MODE

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 FREE REVIEW MODE

– **Amazon Lex** – Providing personalized financial advice through a **ABDULRAH...**

References:

<https://aws.amazon.com/blogs/machine-learning/detect-fraudulent-transactions-using-machine-learning-with-amazon-sagemaker/>

<https://docs.aws.amazon.com/comprehend/latest/dg/tutorial-reviews.html>

<https://aws.amazon.com/lex/financial-services/>

Check out the following Cheat Sheets:

<https://tutorialsdojo.com/amazon-sagemaker/>

<https://tutorialsdojo.com/amazon-comprehend/>

<https://tutorialsdojo.com/amazon-lex/>

15. QUESTION

A data analyst needs to choose the most appropriate visualization methods for various data analysis tasks related to customer purchase patterns.

Match each visualization method with its primary use case in analyzing customer purchase records. Each service should be selected one or more times. (Select THREE.)

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 FREE TIMED MODE

[Free Practice Exam – Timed Mode](#) [1 Topic](#)



 FREE REVIEW MODE

- Bar Chart
- Histogram
- Density plot
- Heatmap

[ABDULRAH...](#)



Examining the frequency of purchases for different items over a specific time period.

Heatmap

(Bar chart)

Understanding the concentration and distribution of transaction values in a continuous range.

Histogram

Identifying patterns in purchasing behavior across various time intervals and item categories.

Bar chart

(Heatmap)

Incorrect

1 / 3 Points

Effective **data visualization** is crucial for assessing the integrity and suitability of data before applying machine learning algorithms. By exploring the visual representation of the data, you can uncover anomalies, biases, and irregularities. Identifying



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

[Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

these issues early on helps ensure that the accuracy of your machine learning models remain robust and reliable.



DATA VISUALIZATION



Bar chart – A bar chart is ideal for visualizing categorical data, such as the frequency of purchases across different product types. Each bar represents a category (e.g., product name), and its length corresponds to the count or frequency, making it simple to compare distinct items over a defined time period.

Histogram – A histogram is designed to display the distribution of continuous numerical data. It groups data into defined intervals (or “bins”) and uses bars to show the number of records that fall within each range. This is especially useful when analyzing transaction values to understand spending patterns or identify unusual concentrations in the data.

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

Heatmap – A heatmap represents two-dimensional data through color intensity, making it useful for detecting patterns across categories and time intervals. For example, in analyzing purchasing behavior, a heatmap can highlight peak purchase times for specific items, revealing underlying trends or seasonal behaviors across the dataset.

Hence, the correct answers are:

- Examining the frequency of purchases for different items over a specific time period: **Bar chart**
- Understanding the concentration and distribution of transaction values in a continuous range: **Histogram**
- Identifying patterns in purchasing behavior across various time intervals and item categories: **Heatmap**

The option that says: **Density plot** is incorrect. A density plot is primarily used to visualize the distribution of continuous data in a smooth, abstract way, but it simply lacks the precision and clarity that histograms offer when analyzing actual transaction frequencies. It is just not ideal for business-oriented scenarios that require exact counts or bin comparisons. Unlike bar charts or heatmaps, a density plot cannot effectively represent categorical data or multi-dimensional



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#) **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

patterns such as item frequency over time. Therefore, while useful in certain contexts, it typically does not provide the actionable insights needed for customer purchase pattern analysis.

References:

<https://aws.amazon.com/what-is/data-visualization/>

<https://docs.aws.amazon.com/quicksight/latest/user/bar-charts.html>

<https://docs.aws.amazon.com/quicksight/latest/user/histogram-charts.html>

Check out this AWS Machine Learning And AI Cheat Sheets:

<https://tutorialsdojo.com/aws-cheat-sheets-aws-machine-learning-and-ai/>

16. QUESTION

A company is looking to use Amazon Bedrock for its AI-driven initiatives. They need a tool that can help create advanced conversational chatbots capable of interacting with customers and providing relevant information from their data systems.



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

Which option best describes a use case that Amazon Bedrock excels at?
ABDULRAH...



- ☐ Developing chatbots that provide relevant information to customers.
- ☐ Predicting future product demand.
- ☐ Predictive maintenance of vehicles.
- ☐ Identifying content within social media images.

Correct

Amazon Bedrock is a service that makes it easy for users to build and scale generative AI applications using pre-trained models from leading AI startups and Amazon itself. One of the key use cases of Amazon Bedrock is to enable the creation of sophisticated chatbots and virtual assistants that can understand and generate natural language, making it highly suitable for customer service applications.

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026


 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)



 **FREE REVIEW MODE**



Amazon Bedrock helps create AI applications with a focus on natural language processing (NLP) and understanding, which are critical components for developing advanced conversational chatbots.

Hence, the correct answer is: **Developing chatbots that provide relevant information to customers.**

The option that says: **Predicting future product demand** is incorrect because Amazon Bedrock's primary focus is on generative AI and NLP applications, not predictive analytics.

The option that says: **Identifying content within social media images** is incorrect this task is better suited for services like Amazon Rekognition, which specializes in image and video analysis. Therefore, this does not satisfy the requirements.

The option that says: **Predictive maintenance of vehicles** is incorrect because this option typically involves analyzing sensor data and applying machine learning models, which is not the primary use case for Amazon Bedrock.

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 FREE TIMED MODE

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 FREE REVIEW MODE

References:

<https://aws.amazon.com/bedrock/> [ABDULRAH...](#)

<https://aws.amazon.com/generative-ai/>

Check out this Amazon Bedrock Cheat Sheet:

<https://tutorialsdojo.com/amazon-bedrock/>

17. QUESTION

A finance company recently moved its on-premises application to AWS, using Amazon EC2 for hosting, Amazon RDS for database management, and Amazon S3 for data storage. The company wants to leverage machine learning to analyze log data from various AWS resources that they are using.

Which AWS service would best support these requirements?

- ☐ AWS Config
- ☐ Amazon CloudWatch Logs Insights
- ☐ Amazon Managed Grafana

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**


☐ AWS CloudTrail

[ABDULRAH...](#)



Correct

Amazon CloudWatch Logs Insights is a powerful tool designed to facilitate in-depth analysis of log data within AWS environments. It enables users to run queries on log data collected from various AWS services and applications in real-time. By utilizing a purpose-built query language, CloudWatch Logs Insights allows for complex analysis and visualization of log data, helping to detect trends, identify performance bottlenecks, and troubleshoot issues efficiently. This capability is particularly valuable for monitoring applications and infrastructure, as it provides actionable insights into system behavior and operational health.



The service integrates seamlessly with Amazon CloudWatch Logs, offering advanced features for querying and visualizing log data. Users can create

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#) FREE TIMED MODE

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 FREE REVIEW MODE

custom dashboards, set up alerts based on specific log patterns, and generate reports to support security and compliance requirements. Users can also use the 'pattern' command, which is backed by machine learning, to automatically detect patterns in log data, assemble similar logs, and summarize vast amounts of log lines into concise, visual formats.

Hence, the correct answer is: **Amazon CloudWatch Logs Insights**.

The option that says: **AWS Config** is incorrect. This service is primarily used for assessing, auditing, and evaluating the configurations of your AWS resources. While it provides valuable insights into configuration drift and compliance status, it does not offer real-time log analysis or querying capabilities for detecting patterns and troubleshooting issues.

The option that says: **AWS CloudTrail** is incorrect because it is just a service that enables governance, compliance, and operational and risk auditing of your AWS account. It records API calls made on your account and provides a history of AWS API calls for your account, including API calls made through the AWS Management Console, AWS



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 FREE TIMED MODE

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 FREE REVIEW MODE

SDKs, command-line tools, and other AWS services. It does not provide real-time log analysis and querying capabilities.

The option that says: **Amazon Managed Grafana** is incorrect. This service is primarily a visualization tool rather than a log analysis service and does not inherently analyze log data or utilize machine learning to identify log patterns.

References:

<https://docs.aws.amazon.com/AmazonCloudWatch/latest/logs/Analyz>

<https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/>

<https://docs.aws.amazon.com/decision-guides/latest/monitoring-on-aws-how-to-choose/monitoring-on-aws-how-to-choose.html>

Check out this Amazon CloudWatch Cheat Sheet:

<https://tutorialsdojo.com/amazon-cloudwatch/>

18. QUESTION

A healthcare analytics team is discussing the lifecycle management of models deployed to production. The team's data scientist is tasked with deciding whether

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)



 **FREE REVIEW MODE**

the deployed machine learning models need to be retrained and updated regularly ABDULRAH...



Which of the following statements accurately reflects the requirements for maintaining machine learning models in production?

- ☐ Production data does not differ from training data over time, so models do not require continuous retraining.
- ☐ Production data is always consistent with training data, making retraining unnecessary.
- ☐ Machine learning models need to be regularly retrained as data changes over time.
- ☐ Once a machine learning model is deployed, it can remain effective indefinitely without updates.

Correct



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026


 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

Machine learning models in production environments are often subject to changes in the underlying data distributions, a phenomenon known as data drift. This can lead to decreased model performance over time if not addressed.



Regular retraining helps models adapt to new data patterns and maintain their accuracy and relevance. Monitoring tools and practices, such as those provided by AWS SageMaker Model Monitor, are crucial for identifying when retraining is necessary.

Hence, the correct answer is: **Machine learning models need to be regularly retrained as data changes over time.**

The option that says: **Production data does not differ from training data over time, so models do not require continuous retraining** is incorrect. Production data can simply change over time due to various factors such as seasonality, trends, or external influences. These changes, known as data

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

drift, can lead to discrepancies between the training data and the production data, necessitating continuous retraining to ensure the model's accuracy and relevance.


The option that says: **Production data is always consistent with training data, making retraining unnecessary** is incorrect because production data is rarely consistent with training data over extended periods. As mentioned above, the various factors cause the production data to diverge from the training data. This inconsistency requires regular monitoring and retraining of the models to ensure they continue to perform well.

The option that says: **Once a machine learning model is deployed, it can remain effective indefinitely without updates** is incorrect because machine learning models must be regularly updated to account for changes in data patterns and maintain their performance. Without updates, models can become outdated and less effective as they fail to adapt to new data characteristics.

References:

<https://docs.aws.amazon.com/sagemaker/latest/dg/model-monitor.html>

<https://docs.aws.amazon.com/sagemaker/latest/dg/whatis.html>



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#) **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

ABDULRAH...
Check out this Amazon SageMaker Sheet:

<https://tutorialsdojo.com/amazon-sagemaker/>



19. QUESTION

A financial company has accumulated a large volume of client feedback, reviews, and comments over the years. The company wants to analyze these data to identify common topics that frequently appear in the feedback to understand client sentiment better and improve their offerings. The specific topics are not predetermined, and the goal is to discover them automatically from the text data.

Which built-in Amazon SageMaker AI algorithm is the most appropriate for this task?

- ☐ Latent Dirichlet Allocation (LDA) Algorithm
- ☐ BlazingText algorithm in Text Classification Mode
- ☐ Text Classification – TensorFlow algorithm

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#) 

 **FREE REVIEW MODE**


☐ Sequence-to-Sequence Algorithm

ABDULRAH...



Incorrect

Amazon SageMaker's Latent Dirichlet Allocation (LDA) algorithm is an unsupervised learning technique designed to represent a collection of documents as a combination of various topics. LDA is primarily used to identify a specified number of topics within a set of text documents. In this context, each document is considered an observation; the words within the documents are the features, and the topics are the categories. Since LDA is an unsupervised method, the topics are not specified up front, and the discovered topics may not necessarily match human categorizations. Instead, LDA learns the topics as a probability distribution over the words in the documents, and each document is characterized as a mixture of these topics.



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#) **FREE TIMED MODE**


[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

In the given scenario, the financial services company aims to analyze client feedback to uncover common themes automatically. Since the themes are not predefined, a topic modeling approach is necessary. The LDA algorithm is designed for this purpose and is a powerful tool for text mining and natural language processing tasks. It allows the company to sift through vast amounts of textual data and discern patterns that might take time to be apparent. This insight can help the company make data-driven decisions to enhance its services and products based on the discovered client sentiments and themes. Thus, the LDA algorithm is the most appropriate choice for this task as it directly addresses the company's requirement to automatically discover latent themes in their text data.

Hence, the correct answer is: **Latent Dirichlet Allocation (LDA) Algorithm.**

The option that says: **Sequence-to-Sequence Algorithm** is incorrect. This algorithm is primarily used for tasks such as language translation or summarization. Sequence-to-sequence models are designed to transform a sequence of elements (such as words in a sentence) into another sequence, and they are not suitable for unsupervised topic discovery.



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

The option that says: **Text Classification – TensorFlow Algorithm** is incorrect. **Text Classification** requires predefined classes or topics. The goal in this scenario is to discover unknown topics, not classify text into existing categories. This algorithm is suitable for tasks where the categories are known and need to be applied to new text data.

The option that says: **BlazingText algorithm in Text Classification mode** is incorrect. BlazingText algorithm in Text Classification mode is a supervised learning approach, which means it requires labeled training data to classify the text. In this scenario, the specific topics of the documents are unknown, making supervised learning methods unsuitable.


References:

<https://docs.aws.amazon.com/sagemaker/latest/dg/lda-how-it-works.html>

<https://docs.aws.amazon.com/sagemaker/latest/dg/lda.html>

<https://docs.aws.amazon.com/sagemaker/latest/dg/algos.html>

Check out this Amazon SageMaker AI Cheat Sheet:



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#) **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

<https://tutorialsdojo.com/amazon-sagemaker/>

[ABDULRAH...](#)



20. QUESTION

A financial institution is working on a machine learning model for predicting loan approval decisions. The institution prioritizes explaining lending decisions to stakeholders in human terms for transparency and fairness.

Which of the following machine learning model statements is true, and which does the institution need to choose to meet the requirement?

☐ A simple model should be chosen because accuracy is more important than explainability in loan approval decisions.

☒ A complex model should be chosen because accuracy is more important than explainability in loan approval decisions.

☐ A complex model should be chosen because explainability is crucial in loan approval decisions.

☐ A simple model should be chosen because explainability is crucial in loan

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA- C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#) **FREE TIMED MODE**[Free Practice Exam – Timed Mode](#) [1 Topic](#) **FREE REVIEW MODE**

Incorrect

Machine learning (ML) is the process of training computers, using math and statistical processes, to find and recognize patterns in data. After patterns are found, ML generates and updates training models to make increasingly accurate predictions and inferences about future outcomes based on historical and new data. For example, ML could help determine the likelihood of a customer purchasing a particular product based on previous purchases by the user or the product's past sales history.

What are simple and complex machine learning models?

Simple and complex ML models differ when balancing a model's accuracy (number of correctly predicted data points) and a model's explainability (how much of the ML system can be explained in "human terms"). The output of a simple ML model may be explainable and produce faster results, but the results may be inaccurate. The output of a complex ML model may be accurate, but the results may be difficult to communicate.



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete


Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**



In the context of developing a machine learning model for predicting loan approvals, a financial institution prioritizing transparency and fairness over maximizing accuracy should opt for a simple model that is easily explainable. Transparency in lending decisions is crucial to ensure fairness and mitigate biases. Simple models such as decision trees or logistic regression provide clear visibility into the factors influencing loan approval, making it easier to explain why a decision was made. While accuracy is important in predicting loan approval decisions, the ability to explain why a particular decision was made is equally critical. Additionally, complex models might offer marginally better predictive accuracy, they often lack transparency, making it difficult to justify decisions to stakeholders and regulators.

Hence, the correct answer is: **A simple model should be chosen because explainability is crucial in loan approval decisions.**

[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#)

 **FREE TIMED MODE**


[Free Practice Exam – Timed Mode](#) [1 Topic](#)

 **FREE REVIEW MODE**

The option that says: **A complex model should be chosen because accuracy is more important than explainability in loan approval decisions** is incorrect because it prioritizes accuracy over explainability, which is the requirement in this scenario. Furthermore, while accuracy is important in predicting loan approvals, it cannot be prioritized over explainability in this context. Financial institutions must ensure their decisions are transparent and fair, which is difficult with complex models like deep neural networks or ensemble methods that lack interpretability. Regulatory bodies often require institutions to provide clear reasons for loan approval or denial, making explainability crucial.

The option that says: **A simple model should be chosen because accuracy is more important than explainability in loan approval decisions** is incorrect. This option correctly identifies the use of a simple model but incorrectly states the reason. The emphasis here should be on explainability, not just accuracy. While simple models can sometimes be less accurate than complex ones, their interpretability makes them more suitable for loan approval decisions where transparency is essential.

The option that says: **A complex model should be chosen because explainability is crucial in loan approval decisions** is incorrect. While it correctly



[Back to Course](#)

FREE AWS Certified Machine Learning Engineer Associate Practice Exams MLA-C01 – Sampler

33% Complete

Last activity on September 5, 2026

 [Your Progress](#) **FREE TIMED MODE**

[Free Practice Exam – Timed Mode](#) [1 Topic](#)



FREE AWS Certified Machine Learning Ei

FREE AWS Certified Machine Learning Ei

 **FREE REVIEW MODE**

states that explainability is crucial, it incorrectly suggests using a complex model. Complex models are typically less interpretable, making it difficult to explain the decision-making process. For ensuring fairness and transparency, simple models that are inherently more interpretable are preferred.

References:

<https://aws.amazon.com/what-is/machine-learning/>

<https://docs.aws.amazon.com/sagemaker/latest/dg/clarify-model-explainability.html>

Check out this AWS Machine Learning And AI Cheat Sheets:

<https://tutorialsdojo.com/aws-cheat-sheets-aws-machine-learning-and-ai/>



ABDUL RAHMAN

ABDUL RAHMAN